

Class 7: Clustering and PCA

A16125573


Clustering

First let's make up some data to cluster so we can get a feel for these methods and how to work with them.

We can use the `rnorm()` function to get random numbers from a normal distribution around a given mean

```
hist(rnorm(5000, mean=3))
```

Histogram of `rnorm(5000, mean = 3)`



```
# mean and sd has a default value, so should worry most about the one that doesn't (n)
```

Let's get 30 points with a mean of 3.

```
c(rnorm(30, mean=3), rnorm(30, mean=-3))
```

```
[1] 3.2084244 3.7054068 3.4705621 4.0661434 2.1897668 4.4194926
[7] 3.1678866 1.9749433 2.1956443 3.9203972 2.7047327 2.7443981
[13] 3.8872643 3.6426726 4.9130119 1.9327703 4.5021159 3.8372291
[19] 4.2150368 0.3765644 3.3919821 0.8314147 3.7281302 3.4500863
[25] 3.6925781 3.3809910 2.2970447 3.1436025 1.2602399 2.2326780
[31] -3.9284313 -2.0047439 -3.7890276 -5.4954627 -4.2042352 -3.4685389
[37] -1.5980628 -3.9160934 -5.3001799 -1.5341460 -2.8632875 -2.9334831
[43] -2.9904012 -3.1754922 -2.1628485 -4.0954682 -2.8243610 -1.8922640
[49] -2.7839677 -3.2195891 -3.0894440 -3.5936179 -2.6009795 -3.7307964
[55] -3.2295156 -3.6557303 -2.7683935 -2.2659478 -2.6431553 -3.5518541
```

```
# cbind()
```

```
rev(c(rnorm(30, mean=3), rnorm(30, mean=-3))) # this reverses
```

```
[1] -4.957870 -2.039953 -2.822553 -3.083814 -3.793521 -4.010679 -1.624195
[8] -3.772316 -2.706706 -1.521207 -2.610659 -4.526746 -3.185522 -3.375776
[15] -4.029854 -3.205724 -2.026053 -3.996235 -1.251944 -1.842881 -2.646989
[22] -3.327836 -2.469409 -4.091714 -2.925673 -2.805249 -2.553231 -3.095059
[29] -3.742598 -3.115158 3.729554 3.195838 4.005090 4.596818 2.562425
[36] 1.526517 4.807164 3.802973 3.453638 3.909998 2.142385 3.294790
[43] 3.250758 4.192773 3.154904 2.261173 2.372802 3.499841 1.658599
[50] 3.725033 2.692871 3.001185 3.191584 3.432434 3.625258 5.123551
[57] 3.204323 6.352056 4.107484 5.263905
```


```
tmp <- c(rnorm(30, mean=3), rnorm(30, mean=-3))
```

```
tmp
```

```
[1] 4.2396827 2.1934619 2.6626342 2.5859659 3.2736025 3.0643913
[7] 0.7411759 1.0499690 3.0059509 2.3151920 1.4878670 0.8115113
[13] 4.3109357 2.4723246 1.5606790 2.9477691 4.2211109 2.5121841
[19] 3.6841331 3.6870870 2.9915633 2.5407447 2.4924749 2.7678030
[25] 2.5044406 2.0448785 1.8782077 4.4486437 3.8326469 2.0627920
[31] -3.3401594 -2.5026304 -2.9051613 -3.4244809 -3.5874678 -2.8765918
```

```
[37] -3.7255310 -2.7810177 -3.3638409 -3.4992866 -1.4970091 -3.2426324
[43] -4.4608747 -3.4696607 -2.9987337 -3.2484622 -2.1292967 -3.3894992
[49] -1.6883479 -4.1248860 -1.3782282 -3.9333013 -2.3230910 -2.4490755
[55] -3.8035898 -3.8542453 -3.4318094 -2.7434786 -2.7294987 -1.6021322
```

```
x <- cbind(x=tmp, y=rev(tmp))
plot(x)
```



K-means clustering.

Very popular clustering method, especially for big datasets, that we can use with the `kmeans()` function in base R.

```
# kmeans() #needs 2 things/inputs

km <- kmeans(x, centers=2)
# centers = # of clusters

km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```
      x      y
1 -3.016801  2.679727
2  2.679727 -3.016801
```

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 47.68179 47.68179
(between_SS / total_SS =  91.1 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"       "
```

```
# we know the answer should be 30, 30
# cluster means = the center points of the clusters
# clustering vector = shows which cluster each point belongs to
# available components = ?
```

```
km$size
```

```
[1] 30 30
```

```
# example data
tmp <- c(rnorm(30, 3), rnorm(30, -3))
x <- data.frame(x=tmp, y=rev(tmp))

#plot(x)

km <- kmeans(x, centers=2, nstart=20)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	3.178960	-3.104514
2	-3.104514	3.178960

Clustering vector:

[illegible]

Within cluster sum of squares by cluster:

```
[1] 55.20742 55.20742
      (between_SS / total_SS =  91.5 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Q. How many points are in each cluster?

km\$size

[1] 30 30

Q. How do we get to cluster membership/assignment?

km\$cluster

[illegible]


Q. How many cluster centers?

km\$centers

	x	y
1	3.178960	-3.104514
2	-3.104514	3.178960


Q. Plot x colored by the kmeans cluster assignment and add cluster centers as blue points

```
# mycols <- km$cluster
plot(x, col=km$cluster)
points(km$centers, col="blue", pch=15, cex=3)
```



Q Let's cluster into 3 groups or same x data and make a plot.

```
km3 <- kmeans(x, centers=3)
plot(x, col=km3$cluster)
```



Hierarchical Clustering

We can use the `hclust()` function for Hierarchical Clustering. Unlike `kmeans()`, where we could just pass in our data as input, we need to give `hclust()` a “distance matrix”

We will use the `dist()` function to start with.

```
d <- dist(x) # a distance matrix
hc <- hclust(d)
hc
```


Call:

```
hclust(d = d)
```

```
Cluster method   : complete
Distance         : euclidean
Number of objects: 60
```

```
plot(hc)
```

Cluster Dendrogram



```
hclust (*, "complete")
```

I can now cut my tree with the `cutree()` to yield a cluster membership vector.

```
grps <- cutree(hc, h=8)
grps
```


```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

You can also tell `cutree()` to cut where it yields “k” groups.

```
cutree(hc, k=2) #cut it into 2 groups
```

[1] 1 2 2 2 2 2 2 2
[39] 2

```
plot(x, col=grps)
```

Principal Component Analysis (PCA)

PCA of UK food

```
url <- "https://tinyurl.com/UK-foods"  
x <- read.csv(url, row.names=1)
```

```
# Complete the following code to find out how many rows and columns are in x?  
nrow(x)
```

```
[1] 17
```

```
ncol(x)
```

```
[1] 4
```

```
dim(x)
```

```
[1] 17 4
```

```
# Preview the first 6 rows
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

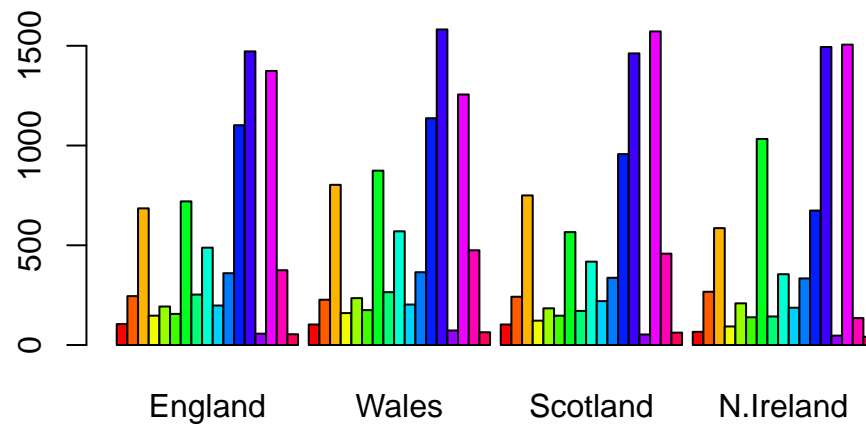
```
#rownames(x) <- x[,1]
#x <- x[,-1]
#head(x)
dim(x)
```

```
[1] 17 4
```

```
# I don't prefer this method because each time you run it, you would drop a column
```

Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

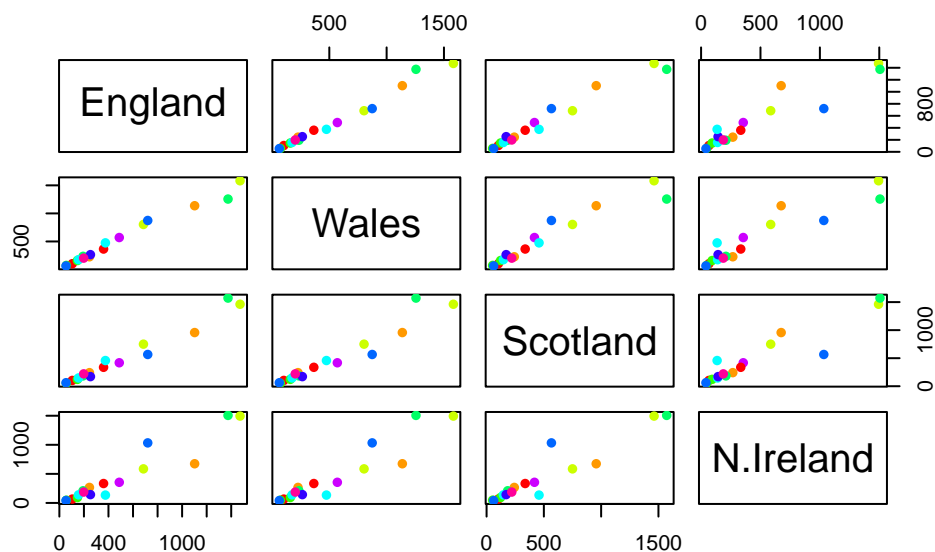
```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



```
# changing beside to F will generate the other bar plot
```

Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch=16)
```



```
# the figure shows
# each colored dot is a food category
# if a points lies on the diagonal, it means that the the x and y axis countries consumed
# points above the diagonal means that the y axis country has consumed more
```

The main PCA function in base R is called `prcomp()` it expects the transposed of our data.

```
# prcomp
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	4.189e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
attributes(pca) # we want x
```

```
$names
```

```
[1] "sdev"      "rotation" "center"    "scale"     "x"


$class
[1] "prcomp"
```

```
pca$x # shows where each country lies on the new axis
```

	PC1	PC2	PC3	PC4
England	-144.99315	2.532999	-105.768945	2.842865e-14
Wales	-240.52915	224.646925	56.475555	7.804382e-13
Scotland	-91.86934	-286.081786	44.415495	-9.614462e-13
N.Ireland	477.39164	58.901862	4.877895	1.448078e-13

```
# Plot PC1 vs PC2
```

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500), col=c("darkorange", "darkgreen", "darkblue", "darkred"))
```



```
# text(pca$x[,1], pca$x[,2], colnames(x))
```